

# GroundedNutriRec Framework

Literature and Concept Notes: Health-Aware and Explainable Food Recommendation Systems Using Retrieval-Augmented LLMs

**Course:** Summer Internship Research Program

**Role:** Data + Preprocessing Lead

**Deliverable:** Foundational Week 1 — System Architecture Notes

**Document Type:** Literature Review and Conceptual Framework

|                      |                                |                   |
|----------------------|--------------------------------|-------------------|
| Recommender Systems  | Food Recommendation            | Health-Aware AI   |
| Explainable AI (XAI) | Retrieval-Augmented Generation | LLM Hallucination |

## TABLE OF CONTENTS

---

### 1. Recommender Systems

- 1.1 Introduction and Architecture
- 1.2 Collaborative Filtering, Matrix Factorization, Neural CF
- 1.3 Content-Based Filtering
- 1.4 Hybrid Systems, Advantages, Disadvantages, Use Cases

### 2. Food Recommendation Systems

- 2.1 Domain-Specific Challenges
- 2.2 Flavor Graphs, Session-Based and Multi-User Models
- 2.3 Advantages, Disadvantages, Use Cases

### 3. Health-Aware Recommendation Systems

- 3.1 Multi-Objective Optimization Framework
- 3.2 Nutritional Scoring, Macronutrient Alignment, Constraints
- 3.3 Advantages, Disadvantages, Use Cases

### 4. Explainable Recommendation Systems

- 4.1 Transparency and Trust Objectives (Tintarev & Masthoff, 2011)
- 4.2 LLM-Powered Explanation Generation in GroundedNutriRec
- 4.3 Advantages, Disadvantages, Use Cases

### 5. Retrieval-Augmented Generation (RAG)

- 5.1 Vector Indexing and ANN Retrieval
- 5.2 Augmented Prompt Construction and Grounded Generation
- 5.3 Advantages, Disadvantages, Use Cases

### 6. LLM Hallucinations in Explanations

- 6.1 Parametric Memory Override and Attention Drift
- 6.2 Claim-Level Faithfulness Verification Pipeline
- 6.3 Faithfulness Score Metric, Advantages, Disadvantages

### 7. Bibliography

- 16 cited references (IEEE / ACM format)

## SECTION 1

# Recommender Systems

### SECTION OVERVIEW

Recommender systems are algorithmic frameworks designed to address information overload by suggesting items that align with a user's latent preferences. They underpin personalization pipelines in e-commerce, streaming platforms, and intelligent food guidance systems including the GroundedNutriRec architecture.

## 1.1 Introduction

Recommender systems map user profiles to item catalogs to output ranked recommendation lists. Formally, let  $U$  be the set of users and  $I$  be the set of items. The objective is to learn a utility function  $R: U \times I \rightarrow S$  that predicts the interaction probability of user  $u$  on item  $i$ , where  $S$  is a totally ordered set (e.g., rating scale  $[0, 5]$  or binary click/no-click). [Adomavicius and Tuzhilin, 2005]

## 1.2 Collaborative Filtering (CF)

Collaborative filtering relies solely on historical interaction patterns between users and items. The foundational assumption is that users who agreed on items in the past will agree in the future [Sarwar et al., 2001]. Two primary CF paradigms exist:

- Matrix Factorization (MF): Learns low-dimensional latent vectors  $p_u, q_i$  in  $\mathbb{R}^d$  for users and items. The predicted rating is formulated as:

Eq. 1 – Matrix Factorization Prediction [Koren et al., 2009]

Predicted Rating:  $\hat{r}(u,i) = \mu + b_u + b_i + p_u^T \cdot q_i$

$\mu$  = global mean rating bias  
 $b_u$  = user-specific bias offset  
 $b_i$  = item-specific bias offset  
 $p_u$  = user latent factor vector (d-dimensional)  
 $q_i$  = item latent factor vector (d-dimensional)

- Deep Collaborative Filtering: Neural Collaborative Filtering (NCF) replaces dot products with multilayer perceptrons. Graph Neural Networks (e.g., LightGCN) capture high-order user-item relationships via node embeddings over bipartite interaction graphs [He et al., 2017].

## 1.3 Content-Based Filtering (CBF)

Content-based filtering recommends items similar to those the user has engaged with, based on item descriptors and user preference vectors mapped to a shared feature space [Lops et al., 2011]. The ranking score is derived via cosine similarity:

**Eq. 2 – Cosine Similarity for Content-Based Scoring**

$$\text{Similarity: } \text{sim}(u, i) = (\theta_u \cdot x_i) / (||\theta_u|| * ||x_i||)$$

$\theta_u$  = user preference vector

$x_i$  = item feature vector

## 1.4 Hybrid Systems

Hybrid systems integrate collaborative and content-based approaches via weighted combination, switching strategies, or feature cascades to simultaneously mitigate cold-start and sparsity limitations.

| Advantages                                                                                                                                                                                                                                                                                           | Disadvantages                                                                                                                                                                                                                                                                                                                                            |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"> <li>Personalizes content at scale, increasing engagement and reducing decision fatigue.</li> <li>Exploits the long tail — exposes niche items unreachable via manual browsing.</li> <li>Simplifies decision-making by prefiltering irrelevant content.</li> </ul> | <ul style="list-style-type: none"> <li>Cold-Start: cannot recommend to new users (user cold-start) or for new items (item cold-start).</li> <li>Data Sparsity: interaction matrices exceed 99% sparsity in real-world deployments.</li> <li>Popularity Bias: models amplify popular item exposure, creating a self-reinforcing feedback loop.</li> </ul> |

## 1.5 Use Cases and Limitations

- E-commerce platforms suggesting related products or automated bundle recommendations.
- Streaming providers dynamically constructing personalized homepages from viewing history.
- Social networks prioritizing news feed items and connection suggestions.

Limitations: Systems cannot recommend items outside the catalog, are vulnerable to shilling attacks that bias ratings, and fail to adapt dynamically to rapid intent shifts without session-aware modeling.

## SECTION 2

# Food Recommendation Systems

### SECTION OVERVIEW

*Food recommendation extends beyond standard user-item matching by incorporating culinary semantics, biochemical flavor constraints, time-structured consumption behaviors, and multi-user aggregation requirements. These domain-specific demands necessitate specialized architectures that generic recommender frameworks do not natively support.*

## 2.1 Domain-Specific Input Complexity

Food recommender systems must process hierarchical recipe inputs combining free-form text, structured ingredient lists, and sequential instructional steps. Three primary NLP processing layers are required [Trattner and Elswiler, 2017]:

- ▣ Title and Description: Free-form natural language requires semantic NLP to extract themes and cuisine types.
- ▣ Ingredient Lists: Named Entity Recognition (NER) normalizes heterogeneous ingredient references (e.g., 'yellow onions', 'chopped onion', 'diced onion' all resolve to a single canonical ingredient node).
- ▣ Instructional Steps: Sequential step parsing models preparation complexity, cooking duration, and technique dependencies.

## 2.2 Flavor Graphs and Session-Based Modeling

Traditional content-based representations fail to capture chemical flavor compound relationships. Advanced architectures construct flavor pairing graphs representing ingredient-level compound compatibility [Ahn et al., 2011]. Additionally, user food consumption is temporally structured around meal schedules and seasonal cycles. Sequential models — including Markov chains and recurrent neural networks — are necessary to prevent recommending heavy dinner recipes during breakfast contexts.

## 2.3 Multi-User Aggregation

Food preparation is frequently a group activity (family meals, social gatherings). Group recommendation algorithms must aggregate individual utility functions to maximize collective satisfaction while enforcing hard constraints such as dietary allergies. This multi-user aggregation problem introduces significant combinatorial complexity into the optimization objective.

| Advantages                                                                                                                                                                                                                                                                                                                                              | Disadvantages                                                                                                                                                                                                                                                                                                                                                                                     |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Facilitates structured meal planning and encourages home cooking by matching available ingredients.</li><li>Reduces food waste by identifying recipes that consume near-expiry inventory items.</li><li>Drives recipe discovery across cuisines by matching flavor profiles across cultural boundaries.</li></ul> | <ul style="list-style-type: none"><li>High model complexity: parsing structured text, raw instruction sequences, and nutritional vectors simultaneously.</li><li>Interaction data is highly sparse; users cook a negligible fraction of the full recipe catalog.</li><li>Offline evaluation metrics (Precision, Recall) do not capture physical variables like ingredient availability.</li></ul> |

## 2.4 Use Cases and Limitations

- Recipe indexing platforms providing personalized search results and meal-kit matching.
- Smart refrigerator interfaces suggesting dishes from live inventory tracking sensors.
- Grocery delivery platforms automating personalized shopping list generation.

Limitations: Cannot verify kitchen tool availability or cooking skill prerequisites. Struggles to model subjective sensory preferences (individual spice tolerances) and cannot dynamically verify ingredient freshness without external sensor integration.

## SECTION 3

# Health-Aware Recommendation Systems

### SECTION OVERVIEW

Health-aware recommenders transition from pure preference satisfaction to multi-objective optimization, balancing taste predictions against nutritional targets and clinical dietary constraints. This shift introduces quantifiable trade-off parameters and hard constraint filters that transform the standard utility function formulation [Ge et al., 2015].

## 3.1 Multi-Objective Utility Formulation

The ranking score for a user-recipe pair  $(u, i)$  is defined as a weighted combination of personalized taste preference and a nutritional healthiness score:

Eq. 3 — Multi-Objective Utility Function [Ge et al., 2015]

$$\text{Utility}(u, i) = \alpha * \text{Pref}(u, i) + (1 - \alpha) * \text{Health}(i)$$

$\text{Pref}(u, i)$  = personalized preference score (from CF or CBF model)

$\text{Health}(i)$  = nutritional healthiness score of recipe  $i$

$\alpha$  = trade-off parameter,  $\alpha$  in  $[0, 1]$

$\alpha \rightarrow 1$  : pure preference optimization (ignores health)

$\alpha \rightarrow 0$  : pure health optimization (ignores taste)

## 3.2 Nutritional Scoring Components

The healthiness score  $\text{Health}(i)$  is computed using three complementary nutritional evaluation layers:

- ❑ FSA/WHO Traffic Light System: Evaluates fat, saturated fat, sugar, and sodium density per 100g, assigning red (high), amber (medium), or green (low) penalty scores [WHO, 2003].
- ❑ Macronutrient Alignment: Measures the deviation of the recipe's protein, carbohydrate, and fat distribution from the user's daily macronutrient targets (e.g., 40% protein / 30% carb / 30% fat for high-performance fitness goals).
- ❑ Dietary Constraint Filters (Hard): If a recipe contains a declared allergen or exceeds diabetic sugar thresholds, its utility score is set unconditionally to zero regardless of the preference score.

Advantages

Disadvantages

|                                                                                                                                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                    |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>□ Promotes population-level health by steering dietary choices away from chronic disease risk factors.</li><li>□ Personalizes clinical dietary therapy – adapts prescribed meal plans to individual taste preferences.</li><li>□ Prevents adverse health events by enforcing allergen and contra-indicated ingredient filters.</li></ul> | <ul style="list-style-type: none"><li>□ Preference Conflict: healthy recipes frequently clash with high-sugar/high-fat taste preferences, reducing CTR.</li><li>□ Data Incompleteness: nutritional scoring requires verified databases and accurate portion sizes.</li><li>□ Dynamic Needs: user nutritional targets change with health status, age, and activity level.</li></ul> |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

### 3.3 Use Cases and Limitations

- Dietary management applications integrated with Electronic Health Records (EHR) for post-discharge patient support.
- Corporate wellness portals suggesting healthy canteen selections to employees.
- Fitness trackers generating meal plans calibrated to real-time calorie expenditure data.

Limitations: Cannot substitute for clinical dietitians in acute pathological conditions. Cannot verify user adherence to recommendations. Struggles with conflicting health objectives (e.g., simultaneous low-sugar diabetic and low-potassium renal dietary requirements).



## SECTION 4

# Explainable Recommendation Systems

### SECTION OVERVIEW

Explainable recommenders provide natural language justifications alongside recommendations, shifting the system from a black-box predictor to a transparent decision-support tool. In GroundedNutriRec, explanations bridge the gap between taste predictions and health objectives, improving dietary adherence through user trust and informed decision-making [Zhang and Chen, 2020].

## 4.1 Explanation Evaluation Objectives

According to Tintarev and Masthoff [2011], recommendation explanations serve seven distinct evaluation objectives that together define the quality of a user-facing explanation system:

| Objective      | Definition                                                                                      |
|----------------|-------------------------------------------------------------------------------------------------|
| Transparency   | Reveals how the system generates recommendations to establish foundational user trust.          |
| Scrutability   | Allows users to identify and correct incorrect model assumptions about their preferences.       |
| Trust          | Increases user confidence in system recommendations through evidence-based justifications.      |
| Persuasiveness | Motivates users to try recommended items by highlighting compelling attributes.                 |
| Efficiency     | Accelerates user decision-making by preemptively addressing the most relevant item attributes.  |
| Satisfaction   | Enhances the overall interaction experience by providing informative and engaging explanations. |
| Effectiveness  | Produces explanations that demonstrably improve decision accuracy and quality outcomes.         |

## 4.2 LLM-Powered Explanations in GroundedNutriRec

GroundedNutriRec employs large language models instructed to generate structured natural language explanations that explicitly link health objectives to user preference history. Example explanation for a hypertensive user:

**Explanation Template — GroundedNutriRec LLM Output**

Example Output: "This dish is recommended because it aligns with your low-sodium dietary goal (142mg/serving) and contains spinach and grilled chicken, which appear in 78% of your previously liked high-protein meals."

| Advantages                                                                                                                                                                                                                                                                                                              | Disadvantages                                                                                                                                                                                                                                                                                                                                                  |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▣ Increases user trust and adoption by making recommendation logic inspectable.</li><li>▣ Empowers users to identify hidden allergens or misaligned assumptions in recommendations.</li><li>▣ Enhances dietary adherence by highlighting specific nutritional benefits.</li></ul> | <ul style="list-style-type: none"><li>▣ High computational cost and latency from generative LLM inference per recommendation.</li><li>▣ Persuasive explanations can convince users to accept poor or clinically inappropriate recommendations.</li><li>▣ Quantitative evaluation of explanation quality is difficult due to subjective user utility.</li></ul> |

## SECTION 5

# Retrieval-Augmented Generation (RAG)

### SECTION OVERVIEW

RAG is a hybrid framework combining dense information retrieval with large language model generation to produce text grounded in external, verifiable documents. In GroundedNutriRec, RAG ensures that all recipe attribute claims in generated explanations are anchored to actual database values rather than the model's parametric memory [Lewis et al., 2020].

## 5.1 Pipeline Architecture

The RAG pipeline in GroundedNutriRec is divided into three sequential phases:

### RAG Pipeline — GroundedNutriRec Three-Phase Architecture

#### Phase 1 — Vector Indexing:

Recipes, ingredient profiles, and USDA nutritional guides are chunked and embedded via SentenceTransformers into d-dimensional vectors. All vectors are indexed in FAISS (Facebook AI Similarity Search).

#### Phase 2 — Retrieval:

Query = (user\_profile + recipe\_id) is embedded.  
FAISS performs Approximate Nearest Neighbor (ANN) search:  
$$\text{sim}(q, d) = (E(q) \cdot E(d)) / (||E(q)|| * ||E(d)||)$$
  
Top-K most similar documents are retrieved.

#### Phase 3 — Augmented Generation:

Retrieved chunks prepended to LLM system prompt as grounding context.  
LLM is instructed: 'Generate explanation ONLY using provided context.'

## 5.2 Cosine Similarity for ANN Retrieval

### Eq. 4 — ANN Cosine Similarity [Lewis et al., 2020]

Retrieval Score: 
$$\text{sim}(q, d) = (E(q) \cdot E(d)) / (||E(q)|| * ||E(d)||)$$

$E(q)$  = dense embedding vector of the user query  
 $E(d)$  = dense embedding vector of a stored document chunk  
 $\text{sim}$  = cosine similarity in range [-1, +1]

| Advantages                                                                                                                                                                                                                                                                                                                                                              | Disadvantages                                                                                                                                                                                                                                                                                                                                            |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>□ Dramatically reduces LLM hallucinations by grounding output in verified source documents.</li><li>□ Knowledge updates are instant — editing database files automatically updates model outputs without retraining.</li><li>□ Provides full output traceability by linking generated text to indexed source documents.</li></ul> | <ul style="list-style-type: none"><li>□ Increased end-to-end latency from vector database retrieval prior to LLM inference.</li><li>□ Retrieval quality is highly sensitive to document chunking strategies and embedding model selection.</li><li>□ Higher storage footprint for maintaining dual-format vector indices and raw text corpora.</li></ul> |

### 5.3 Use Cases and Limitations

- Enterprise knowledge search engines providing grounded, source-attributed factual answers.
- Clinical decision-support assistants summarizing medical research for diagnostic support.
- Explainable food recommenders generating nutritionally justified recipe suggestions.

Limitations: Performance degrades if the vector index contains contradictory or low-quality documents. Multi-hop reasoning across disconnected document sources requires agent-level orchestration beyond standard RAG. All generation is bounded by the LLM context window length.

## SECTION 6

# LLM Hallucinations in Explanations

### SECTION OVERVIEW

LLM hallucinations occur when a generative model produces text that is grammatically coherent but factually incorrect, fabricated, or unsupported by the provided context [Ji et al., 2023]. In the context of explainable food recommendation, hallucinations pose direct safety and clinical credibility risks, particularly when explanations make false allergen or nutritional claims.

## 6.1 Root Causes of Hallucination

Two primary mechanisms drive LLM hallucinations in RAG-based explanation pipelines:

- ▣ Parametric Memory Overriding: The model prioritizes associations learned during pretraining over the structured context provided in the retrieval-augmented prompt, generating plausible but unverified claims.
- ▣ Attention Drift: In long input sequences, attention weights decay over distal tokens, causing the model to generate text anchored in local linguistic patterns rather than the full system prompt context.

## 6.2 Claim-Level Faithfulness Verification Pipeline

GroundedNutriRec implements a three-stage post-generation verification pipeline using a Natural Language Inference (NLI) model to assess each generated explanation:

### GroundedNutriRec – Claim-Level Faithfulness Verification Pipeline

#### Stage 1 – Claim Extraction:

Parse generated explanation into atomic fact units.

Example: 'contains 30g protein', 'prepared in under 15 minutes'

#### Stage 2 – NLI Entailment Classification (per claim):

Entailment : Claim is fully supported by retrieved document context.

Contradiction: Claim directly contradicts the document context.

Neutral : Claim is unverified by context (hallucination candidate).

#### Stage 3 – Faithfulness Score Computation:

$\text{Faithfulness} = (\text{Supported Claims}) / (\text{Total Claims})$

If score < threshold: REJECT explanation -> trigger regeneration.

## 6.3 Faithfulness Score Metric

Eq. 5 — Faithfulness Score Definition [Ji et al., 2023]

$$\text{Faithfulness Score} = |\{\text{claims: NLI\_label} = \text{Entailment}\}| / |\{\text{all\_claims}\}|$$

Score in [0, 1]

Score = 1.0 : all claims are fully document-grounded (ideal)

Score < tau : explanation rejected (tau = configured threshold)

| Advantages                                                                                                                                                                                                                                                                                                                                                   | Disadvantages                                                                                                                                                                                                                                                                                                                                                                  |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"> <li>▣ Provides quantifiable hallucination rate benchmarks for LLM model comparison and selection.</li> <li>▣ Prevents direct medical and dietary safety risks from false allergen or nutrition claims.</li> <li>▣ Atomic claim verification enables targeted explanation repair rather than full regeneration.</li> </ul> | <ul style="list-style-type: none"> <li>▣ High additional latency: NLI inference runs on every claim extracted from every explanation.</li> <li>▣ NLI models themselves introduce classification errors and are subject to their own training biases.</li> <li>▣ Cannot detect implicit misinformation that does not explicitly contradict documented context facts.</li> </ul> |

## 6.4 Use Cases and Limitations

- ▣ Safety guardrails for generative models deployed in high-stakes finance, legal, and medical systems.
- ▣ Summarization pipeline auditing to verify factual accuracy of document condensations.
- ▣ RAG explanation output validation before user-facing display in GroundedNutriRec.

Limitations: Faithfulness verification is a filter layer, not a prevention mechanism — it cannot eliminate hallucinations at the source. Classifying complex, compound, or conditional claims (e.g., comparative nutritional assertions) exceeds the capability of standard NLI models. Verification coverage is bounded by the accuracy and completeness of the ground-truth document database.

---

## SECTION 7

# Bibliography

All references are formatted in accordance with ACM/IEEE citation conventions. Entries are listed in the order of first appearance in the document.

- [1] Adomavicius, G. and Tuzhilin, A., 2005. Toward the next generation of recommender systems: A survey of the state-of-the-art and possible extensions. *IEEE Transactions on Knowledge and Data Engineering*, 17(6), pp.734-749.
- [2] Ahn, Y.Y., Ahnert, S.E., Bagrow, J.P. and Barabasi, A.L., 2011. Flavor network and the principles of food pairing. *Scientific Reports*, 1, p.196.
- [3] Elswailer, D., Ludwig, S., Klempis, M. and Trattner, C., 2017. Promoting healthy food choices via food recommendation. *Proc. ACM CSCW*, pp.2173-2185.
- [4] Freyne, J. and Berkovsky, S., 2010. Recommending food: Recipes, ingredients and ingredient groups. *Proc. ACM RecSys*, pp.245-248.
- [5] Ge, M., Elswailer, D. and Trattner, C., 2015. Health-aware food recommendation in social networks. *IEEE Access*, 3, pp.2181-2191.
- [6] He, X., Liao, L., Zhang, H., Nie, L., Hu, X. and Chua, T.S., 2017. Neural collaborative filtering. *Proc. 26th International Conference on World Wide Web*, pp.173-182.
- [7] Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y.J., Madotto, A. and Fung, P., 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), pp.1-38.
- [8] Konstan, J.A. and Riedl, J., 2012. Recommender systems: From algorithms to user experience. *User Modeling and User-Adapted Interaction*, 22(1), pp.101-123.
- [9] Koren, Y., Bell, R. and Volinsky, C., 2009. Matrix factorization techniques for recommender systems. *Computer*, 42(8), pp.30-37.
- [10] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Kuttler, H., Lewis, M., Yih, W.T., Rocktaschel, T. and Riedel, S., 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, pp.9459-9474.
- [11] Lops, P., De Gemmis, M. and Semeraro, G., 2011. Content-based recommender systems: State of the art and trends. *Recommender Systems Handbook*, pp.73-105.
- [12] Sarwar, B., Karypis, G., Konstan, J. and Riedl, J., 2001. Item-based collaborative filtering recommendation algorithms. *Proc. 10th International Conference on World Wide Web*, pp.285-295.
- [13] Tintarev, N. and Masthoff, J., 2011. Designing and evaluating explanations for recommender systems. *Recommender Systems Handbook*, pp.479-510.
- [14] Trattner, C. and Elswailer, D., 2017. Study on recipes on the web: Analysis of ingredients, preparation time, and nutritional values. *Proc. ACM Hypertext*, pp.25-34.
- [15] World Health Organization (WHO), 2003. Diet, Nutrition and the Prevention of Chronic Diseases. WHO Technical Report Series No. 916. Geneva: WHO.
- [16] Zhang, Y. and Chen, X., 2020. Explainable recommendation: A survey and new perspectives. *ACM Transactions on Information Systems*, 38(4), pp.1-101.